

Vector Processors

Slide courtesy: S. R. Sarangi

Slides adapted by: S. Mittal

Vector Processors

- * A **vector instruction** operates on arrays of data
 - * Example : There are **vector instructions** to add or **multiply** two **arrays** of data, and produce an array as output
 - * **Advantage** : Can be used to perform all kinds of array, matrix, and linear algebra **operations**. These **operations** form the core of many scientific programs, high intensity graphics, and data analytics applications.


Background

- * **Vector processors** were traditionally used in supercomputers
- * **Vector instructions** gradually found their way into **mainstream** processors
 - * MMX, SSE1, SSE2, SSE3, SSE4, and AVX instruction sets for x86 processors
 - * AMD 3D Now Instruction Set

Software Interface

- * Let us define a **vector register**
 - * Example : 128 bit registers in the **MMX instruction set**
→ XMM0 ... XMM15
 - * Can hold 4 floating point values, or 8 2-byte short integers
 - * **Addition** of **vector registers** is equivalent to pairwise addition of each of the individual **elements**.
 - * The **result** is saved in a **vector register** of the same size.

Example of Vector Addition



Let us define 8 128 bit vector registers in SimpleRisc. vr0 ... vr7

Loading Vector Registers

- * There are two **options** :
 - * **Option 1** : We assume that the data elements are stored in **contiguous locations**
 - * Let us define the v.ld **instruction** that uses this **assumption**.

Instruction	Semantics
v.ld vr1, 12[r1]	$vr1 \leftarrow ([r1+12], [r1+16], [r1+20], [r1+24])$

- * **Option 2**: Assume that the elements are not saved in contiguous locations.

Scatter Gather Operation

- * The data is **scattered** in memory
 - * The **load operation** needs to **gather** the data and save it in a **vector register**.
- * Let us define a **scatter gather** version of the **load instruction** \rightarrow v.sg.ld
 - * It uses another **vector register** that contains the addresses of each of the elements.

Instruction	Semantics
v.sg.ld vr1, vr2	$vr1 \leftarrow ([vr2[0]], [vr2[1]], [vr2[2]], [vr2[3]])$

Vector Store Operation

- * We can similarly define two vector store operations

Instruction	Semantics
v.sg.st vr1, vr2	$\begin{aligned} [vr2[0]] &\leftarrow vr1[0] \\ [vr2[1]] &\leftarrow vr1[1] \\ [vr2[2]] &\leftarrow vr1[2] \\ [vr2[3]] &\leftarrow vr1[3] \end{aligned}$

Instruction	Semantics
v.st vr1, 12[r1]	$\begin{aligned} [r1+12] &\leftarrow vr1[0] \\ [r1+16] &\leftarrow vr1[1] \\ [r1+20] &\leftarrow vr1[2] \\ [r1+24] &\leftarrow vr1[3] \end{aligned}$

Vector Operations

- * We can now define **custom operations** on vector registers
 - * v.add → Adds two vector registers
 - * v.mul → Multiplies two vector registers
 - * We can even have operations that have a **vector operand** and a **scalar operand** → Multiply a **vector** with a **scalar**.

Predicated Instructions

- * Suppose we want to run the following **code snippet** on each element of a **vector register**
 - * $\text{if}(x < 10) \ x = x + 10 ;$
- * Let the **input vector register** be vr1
- * We first do a **vector comparison** :
 - * `v.cmp vr1, 10`
 - * It saves the **results** of the comparison in the **v.flags** register (**vector** form of the **flags** register)

Predicated Instructions - II

- * If a condition is true, then the predicated instruction gets evaluated
 - * Otherwise, it is replaced with a nop.
- * Consider a scalar predicated instruction (in the ARM ISA)
 - * addeq r1, r2, r3
 - * $r1 = r2 + r3$ (if the previous comparison resulted in an equality)

Predicated Instructions - III

- * Let us now define a vector form of the **predicated instruction**
 - * For example : `v.<p>.add` (<p> is the **predicate**)
 - * It is a regular **add** instruction for the elements in which the **predicate** is **true**.
 - * For the rest of the **elements**, the instruction becomes a **nop**
- * Example of **predicates** :
 - * It (less than) , gt (greater than), eq (equality)

Predicated Instructions - IV

- * Implementation of our function :

- * if ($x < 10$) $x = x + 20$

```
v.cmp vr1, 10
```

```
v.lt.add vr1, vr1, 20
```

Adds 20 to every element of *vr1*

Design of a Vector Processor

* Salient Points

- * We have a **vector register file** and a **scalar register file**
- * There are scalar and vector functional units
- * Unless we are converting a vector to a scalar or vice versa, we in **general** do not **forward** values between **vector** and **scalar** instructions
- * The **memory** unit needs support for **regular** operations, **vector** operations, and possibly **scatter-gather** operations.